



Course: EE321 Intro to AI

Complex Engineering Problem

Name: Ata Ul Musawir

Reg.: 2023135

IMPORTANT: This is an INDIVIDUAL project

You must run the code yourself.

The exam will ask specific questions about YOUR model's behavior.
If you copy from the internet, you will likely to fail the viva/exam.

Your Must

Build an appropriate AI model that can tell whether a picture contains a **hot dog** or **not a hot dog**.

Learning Objectives

By the end of this project, you will be able to:

- Load and explore image data
- Build an ANN and a CNN in TensorFlow/Keras
- Identify and fix overfitting using Dropout and Data Augmentation
- Plot and interpret training curves
- Test your model on real-world images

Dataset

You will use the **Hot Dog / Not Hot Dog** dataset from Kaggle:

<https://www.kaggle.com/dansbecker/hot-dog-not-hot-dog>

What You Need

- A Google account (for Google Colab)
- The dataset downloaded to your computer (upload to Colab)
- **Submit your overleaf generated PDF by May 07, 2026 by 12:00 PM**

Grading Rubric

Task	Points
Step 2: Data Exploration (answers)	5
Step 3: Visualization (screenshot)	5
Step 5: ANN Model (accuracy + analysis)	10
Step 6: CNN Model (accuracy + comparison)	15
Step 7: Overfitting Fix (before/after plots)	20
Step 8: Real Image Test (screenshot)	10
Step 9: Answers to all questions	20
Step 10: YOLO Extension (extra credit)	5
Total	85 + 5 EC

Step 1: Setup Google Colab

Follow these instructions carefully:

How to Run This Project in Google Colab

1. Go to <https://colab.research.google.com>
2. Click **File** → **New Notebook**
3. Rename it: `HotDog_Classifier_YourName.ipynb`
4. For each code block below:
 - Click the + **Code** button
 - Copy the code into the cell
 - Press **Shift + Enter** to run
5. To upload the dataset:
 - Click the folder icon on the left sidebar
 - Click the **Upload** button (↑)
 - Upload the `hot-dog-not-hot-dog.zip` file
 - Run: `!unzip hot-dog-not-hot-dog.zip`

Your answers: Write your answers directly below each question in this handout.

Step 2: Import Libraries

Run this code in Colab:

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import os
import random

print("Libraries loaded successfully!")
print("TensorFlow version:", tf.__version__)
```

1. What TensorFlow version are you using?

Answer: The TensorFlow version being used is **2.20.0**.

Step 3: Load and Explore the Data

Run this code to see how many images you have:

```
train_path = "/content/seefood/train"
test_path = "/content/seefood/test"

# Count training images
hot_dog_count = len(os.listdir(os.path.join(train_path, "hot_dog")))
not_hot_dog_count = len(os.listdir(os.path.join(train_path, "not_hot_dog")))

print(f"Hot dog images: {hot_dog_count}")
print(f"Not hot dog images: {not_hot_dog_count}")

# Count test images
test_hot = len(os.listdir(os.path.join(test_path, "hot_dog")))
test_not_hot = len(os.listdir(os.path.join(test_path, "not_hot_dog")))

print(f"Test hot dog images: {test_hot}")
print(f"Test not hot dog images: {test_not_hot}")
```

1. How many training images are in each class?

Answer: There are **249 hot dog** images and **249 not hot dog** images in the training set. The test set contains 250 hot dog and 250 not hot dog images.

2. Is the dataset balanced? Why does that matter?

Answer: Yes, the dataset is perfectly balanced — both classes have exactly 249 training images each. This matters because an imbalanced dataset causes the model to become biased toward the majority class, giving misleadingly high accuracy while performing poorly on the minority class.

Step 4: Visualize the Data

```
plt.figure(figsize=(12, 6))

for i in range(6):
    class_name = random.choice(["hot_dog", "not_hot_dog"])
    class_path = os.path.join(train_path, class_name)
    img_name = random.choice(os.listdir(class_path))
    img_path = os.path.join(class_path, img_name)

    img = plt.imread(img_path)
    plt.subplot(2, 3, i + 1)
    plt.imshow(img)
    plt.title(class_name)
    plt.axis('off')

plt.tight_layout()
plt.show()

# Print image shape
print("Image shape:", img.shape)
```

1. Paste a screenshot of your 6 random images here:



Figure 1: 6 random training images from the dataset

2. What is the shape (height, width, channels) of a typical image? (Use `img.shape`)

Answer: The shape of a typical image is **(396, 512, 3)** — height 396 pixels, width 512 pixels, and 3 color channels (RGB).

Step 5: Prepare the Data Generator

```
IMG_HEIGHT = 64
IMG_WIDTH = 64
BATCH_SIZE = 32

train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2
)

train_generator = train_datagen.flow_from_directory(
    train_path,
    target_size=(IMG_HEIGHT, IMG_WIDTH),
    batch_size=BATCH_SIZE,
    class_mode='binary',
    subset='training'
)

validation_generator = train_datagen.flow_from_directory(
    train_path,
    target_size=(IMG_HEIGHT, IMG_WIDTH),
    batch_size=BATCH_SIZE,
    class_mode='binary',
    subset='validation'
)

# Print class labels
print(train_generator.class_indices)
```

1. What does class label 0 represent? What does 1 represent?

Answer: Class label **0** represents **hot_dog** and class label **1** represents **not_hot_dog**, as confirmed by the output: `{'hot_dog': 0, 'not_hot_dog': 1}`. The generator found 400 training images and 98 validation images.

Step 6: Build a Simple ANN (It Will Fail)

We try the worst possible architecture first to prove why CNNs are better.

```
ann_model = keras.Sequential([
    layers.Flatten(input_shape=(IMG_HEIGHT, IMG_WIDTH, 3)),
    layers.Dense(128, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

ann_model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

ann_model.summary()

history_ann = ann_model.fit(
    train_generator,
    validation_data=validation_generator,
    epochs=10
)

# Plot
plt.plot(history_ann.history['accuracy'], label='Train Accuracy')
plt.plot(history_ann.history['val_accuracy'], label='Validation
Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('ANN Training vs Validation Accuracy')
plt.grid(True)
plt.show()
```

1. What was the final validation accuracy of the ANN?

Answer: The final validation accuracy of the ANN after 10 epochs was **58.16%** (val_accuracy = 0.5816). Training accuracy reached 81.00%.

2. Why does the ANN perform poorly on image data? (One sentence)

Answer: The ANN flattens the 2D image into a 1D vector of 12,288 values, destroying all spatial structure, so it cannot detect edges, textures, or local patterns the way a CNN can.

3. Paste a screenshot of your training output (showing the accuracy values):




Figure 2: ANN training output showing accuracy per epoch

Step 7: Build a CNN

Now we use Convolutional and Pooling layers.

```
cnn_model = keras.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(
        IMG_HEIGHT, IMG_WIDTH, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

cnn_model.compile(optimizer='adam',
                  loss='binary_crossentropy',
```



```

        metrics=['accuracy'])

history_cnn = cnn_model.fit(
    train_generator,
    validation_data=validation_generator,
    epochs=15
)

# ANN vs CNN plot
plt.plot(history_ann.history['val_accuracy'], label='ANN Validation',
         marker='o')
plt.plot(history_cnn.history['val_accuracy'], label='CNN Validation',
         marker='s')
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.legend()
plt.title('ANN vs CNN Validation Accuracy')
plt.grid(True)
plt.show()

```

1. What was the final validation accuracy of the CNN?

Answer: The final validation accuracy of the CNN after 15 epochs was **58.16%**. The best validation accuracy reached was **64.29%** at Epoch 12. Training accuracy reached 95.25%, indicating significant overfitting.

2. Compare ANN vs CNN. Which one performed better and why?

Answer: The CNN performed better than the ANN, reaching a best validation accuracy of 64.29% compared to the ANN's best of 57.14%. The CNN uses convolutional filters to detect local spatial features such as edges and textures while preserving spatial relationships that the ANN destroys by flattening. However, both models show overfitting — the CNN's training accuracy (95.25%) far exceeds its validation accuracy (58.16%).

3. Paste the accuracy plot comparing ANN and CNN: (Use the code below to generate the plot)

```

plt.plot(history_ann.history['val_accuracy'], label='ANN
Validation', marker='o')
plt.plot(history_cnn.history['val_accuracy'], label='CNN
Validation', marker='s')
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.legend()
plt.title('ANN vs CNN Validation Accuracy')
plt.grid(True)
plt.show()

```




Figure 3: ANN vs CNN Validation Accuracy comparison

Step 8: Fix Overfitting (Most Important)

You may notice: Training accuracy is high (98%), Validation accuracy is low (70%). This is OVERFITTING. The model memorized the training images but doesn't generalize.

Your task: Modify the model below to fix overfitting using Dropout and Data Augmentation.

8.1 Fix 1: Add Dropout

```
from tensorflow.keras import regularizers

# Augmented data generator
augmented_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
    validation_split=0.2
)

train_augmented = augmented_datagen.flow_from_directory(
```

```

train_path,
target_size=(IMG_HEIGHT, IMG_WIDTH),
batch_size=BATCH_SIZE,
class_mode='binary',
subset='training'
)

# Final model with Dropout + Augmentation
final_model = keras.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(
        IMG_HEIGHT, IMG_WIDTH, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.25),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.25),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(1, activation='sigmoid')
])

final_model.compile(optimizer='adam',
                    loss='binary_crossentropy',
                    metrics=['accuracy'])

history_final = final_model.fit(
    train_augmented,
    validation_data=validation_generator,
    epochs=15
)

# Before vs After plot
plt.plot(history_cnn.history['val_accuracy'], label='CNN (Overfit)',
         marker='o')
plt.plot(history_final.history['val_accuracy'], label='CNN + Dropout +
Augment', marker='s')
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.legend()
plt.title('Improving Generalization')
plt.grid(True)
plt.show()

```

8.2 Fix 2: Add Data Augmentation

Modify the ImageDataGenerator to create more variations:

```
augmented_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
    validation_split=0.2
)

train_augmented = augmented_datagen.flow_from_directory(
    train_path,
    target_size=(IMG_HEIGHT, IMG_WIDTH),
    batch_size=BATCH_SIZE,
    class_mode='binary',
    subset='training'
)

# Then retrain the improved model with augmented_datagen
final_model = keras.Sequential([
    # Same architecture as improved_model
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(
        IMG_HEIGHT, IMG_WIDTH, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.25),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.25),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(1, activation='sigmoid')
])

final_model.compile(optimizer='adam', loss='binary_crossentropy',
    metrics=['accuracy'])

history_final = final_model.fit(train_augmented, validation_data=
    validation_generator, epochs=15)
```

1. Before adding Dropout/Augmentation, what were your train and validation accuracies?

Train: **95%** Validation: **58%**

2. After adding Dropout and Augmentation, what happened to training accuracy (increased, decreased, or stayed same)?

Training accuracy **decreased** — Dropout randomly disables neurons during training, preventing memorization.

Train: **64%**

3. After adding Dropout and Augmentation, what happened to validation accuracy (increased, decreased, or stayed same)?

Validation accuracy **increased** — best validation accuracy improved to 65.31% at Epoch 13, with a smaller gap between train and validation accuracy showing better generalization.

Validation: **65%**

Question 8.3

Paste the before/after validation accuracy plot here:

```
plt.plot(history_cnn.history['val_accuracy'], label='CNN (Overfit)',
         marker='o')
plt.plot(history_final.history['val_accuracy'], label='CNN + Dropout +
         Augment', marker='s')
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.legend()
plt.title('Improving Generalization')
plt.grid(True)
plt.show()
```




Figure 4: Before vs After: CNN with and without Dropout and Augmentation

Step 9: Test on Real Images

Find a hot dog image online (or take a photo) and test your model:

```
from tensorflow.keras.preprocessing import image
import numpy as np
import matplotlib.pyplot as plt

def predict_local_image(img_path):
    img = image.load_img(img_path, target_size=(IMG_HEIGHT,
        IMG_WIDTH))
    img_array = image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0) / 255.0

    prediction = final_model.predict(img_array)[0][0]

    if prediction > 0.5:
        print(f"HOT DOG! (confidence: {prediction:.2f})")
    else:
        print(f"NOT A HOT DOG! (confidence: {1 - prediction:.2f})")

    plt.imshow(img)
    plt.axis('off')
```

```
plt.show()

# Use a test from external image
predict_local_image("/content/hotdog_test.jpg")
```

1. Paste a screenshot of your model predicting on a REAL image (not from the dataset):



Figure 5: Model prediction output: HOT DOG! (confidence: 0.78)

2. Did your model ever make a funny mistake (e.g., calling a hamburger a hot dog)? What did it confuse?

Answer: The model correctly identified the test hot dog image with **78% confidence**. Since the dataset is small (249 training images per class), the model may confuse visually similar foods such as sausages or corn dogs as hot dogs, since they share similar elongated shapes and colors.

Step 10: Optional Stretch Goal

YOLO Object Detection (Extra Credit)

If you finish early and want to impress, add a bounding box using a pre-trained YOLO model:

```
!pip install ultralytics -q

from ultralytics import YOLO

yolo_model = YOLO("yolov8n.pt")

# Use local image instead of URL
results = yolo_model("/content/hotdog_test.jpg")

results[0].show()
```

Extra Credit Question

Does YOLO correctly detect a "hot dog" as a class? What label does it give?

Answer: Yes, YOLO correctly detected **2 hot dogs** in the image and labeled them as "hot dog" — one of the 80 object classes in the COCO dataset that YOLOv8 was pre-trained on. Detection completed in 182.2ms.




Figure 6: YOLOv8 detection: 2 hot dogs detected with bounding boxes

Project Submission Links:**Google Colab Notebook:**

<https://colab.research.google.com/drive/1ArVJo-HwHlP2PqZYXWSa89QvvpUBcT4r?usp=sharing>

GitHub Repository:

<https://github.com/ataulmusawir/Introduction-to-AI-Assignment-2>

Final Submission Checklist

Step 2: TensorFlow version answer

Step 3: Image counts and balance explanation

Step 4: Screenshot of 6 images + shape answer

Step 6: ANN accuracy + explanation + screenshot

Step 7: CNN accuracy + comparison + plot screenshot

Step 8: Overfitting before/after numbers + plot screenshot

Step 9: Real image prediction screenshot + mistake explanation

Step 10: (Extra credit) YOLO results screenshot

Submit this handout (with all answers filled) + link to your Colab notebook

Remember

The exam will ask you specific questions about YOUR project.

If you run the code yourself, you will know the answers easily.

If you copy from someone else, you will fail the exam.

Good luck and have fun!